



Jung Sungwoo

Posted on 2025년 12월 30일



1



1

# World-Centric Agent Architecture: Why Your AI Agent Keeps Failing (And It's Not the Model's Fault)

#agents #architecture #llm #ai

## World-Centric Agent Architecture Deterministic Runtime and Hierarchical Intelligence Orchestration

### TL;DR

Your AI agent's failures aren't because the model isn't smart enough. They're because **the World was never designed**.

This post introduces **World-Centric Architecture**—a paradigm where:

- The World is explicit, immutable, and verifiable

- Intelligence proposes changes but never executes them directly
- System stability is independent of model size
- Every value can answer "Why?"

If you're expecting another agent framework tutorial, this is not it. This post is about why agent systems fail structurally—even with strong models.

## The Problem Nobody's Talking About

LLM-based agents have made remarkable progress. Yet they keep failing in predictable ways:

Failure Mode	What Happens
<b>Impossible Action Loops</b>	Agent repeatedly tries actions that can't work
<b>State Opacity</b>	No one knows when or why state changed
<b>Irreproducibility</b>	Same input, different results
<b>Unrecoverable Failures</b>	Can't rollback, can't try alternatives
<b>Unexplainability</b>	"Why did it do that?" → 🤷‍♂️

Sound familiar?

## The Misdiagnosis

The industry's response has been consistent:

"The model isn't smart enough. Let's use a bigger one."

Problem → "Model too dumb"

- Bigger model / Longer reasoning / Complex planning
- Higher cost / More complexity / Problem persists

We've been treating symptoms, not the disease.



+


...

## The Real Problem

Here's the uncomfortable truth:

**These failures aren't intelligence failures. They're world modeling failures.**

In most agent architectures, the "world" exists only inside the model's head:




The model must:

1. Understand the world (from context)
2. Know what's possible (by reasoning)
3. Pick the right action (while hallucinating)
4. Track state changes (in its attention weights)
5. Explain failures (good luck)

That's too much responsibility for a non-deterministic text predictor.

## The Inversion: World-Centric Design

What if we flipped the script?





Intelligence  
(Proposer)



Intelligence is OUTSIDE. World is explicit.



## Key shifts:

- World exists **outside** the model
- State is **explicit and immutable**
- Actions have **verifiable availability**
- Intelligence **proposes**, doesn't execute
- Everything can **explain itself**

## The Constitution


World-Centric Architecture follows six principles:

### THE CONSTITUTION

1. COMPILER defines possibility.  
What CAN happen is defined at design time.
2. CORE computes truth.  
What IS true is computed deterministically.
3. ACTOR proposes change.  
Intelligence suggests, never executes.
4. AUTHORITY judges proposals.  
Independent verification before any change.
5. ORCHESTRATOR manages worlds.  
Multiple world branches for exploration.
6. PROJECTION transforms I/O.  
Presentation is separate from meaning.

# The Six Pillars

Here's how these principles become architecture:



Pillar	Role	Owns State?	Makes Decisions?
<b>Core</b>	Truth Engine	<input checked="" type="checkbox"/> Snapshot	<span style="color: red;">X</span>
<b>Actor</b>	Change Proposer	<span style="color: red;">X</span>	<span style="color: red;">X</span>
<b>Authority</b>	Proposal Judge	<span style="color: red;">X</span>	<input checked="" type="checkbox"/>
<b>Projection</b>	I/O Transformer	<span style="color: red;">X</span>	<span style="color: red;">X</span>
<b>Orchestrator</b>	World Manager	<input checked="" type="checkbox"/> Runtimes	<span style="color: red;">X</span>
<b>Compiler</b>	Possibility Definer	<input checked="" type="checkbox"/> Schema	<span style="color: red;">X</span>



## Core: The Deterministic Truth Engine

Core is the foundation—a seven-layer engine that computes truth without side effects:

Layer 7: EXPLAIN  
"Why this value?" – Structural interpretation

Layer 6: PATCH / APPLY  
The ONLY way to change truth

Layer 5: EFFECT RUNTIME  
External execution via handlers

Layer 4: ACTION  
Availability gate + Effect declaration

Layer 3: SNAPSHOT  
Immutable state at a point in time

Layer 2: DAG  
Dependency graph for incremental recompute

Layer 1: EXPRESSION  
Pure computation (no side-effects)

## Snapshots Are Immutable

```
type Snapshot<T> = {
  readonly data: T;
  readonly computed: Record<string, unknown>;
  readonly validity: ConstraintResult;
  readonly version: number;
};

// WRONG – Direct mutation
snapshot.data.user.email = 'new@example.com';

// RIGHT – Patch/Apply produces NEW snapshot
const newSnapshot = applyPatches(snapshot, [
  set('user.email', 'new@example.com')
]);
```

## Actions Gate on Availability



```
const SubmitAction = defineAction(AppState, () => ({
  // Availability is a COMPUTED FACT
  availability: get('user.computed.isValid'),

  // Effect only runs IF available
  effect: effect('api.submit', {
    email: get('user.email'),
  }),
}));
```

**Impossible actions never reach execution.** The system prevents them structurally, not through model reasoning.

## Everything Explains Itself

```
const result = explainAvailability(SubmitAction, snapshot);

// Returns structured explanation:
{
  available: false,
  tree: {
    operator: 'and',
    contribution: 'false',
    children: [
      { path: 'user.computed.emailOk', value: false }, // ←
      { path: 'user.computed.passwordOk', value: true },
    ]
  }
}
```

No guessing. No "maybe the email was empty?" Just structural facts.

## Hierarchical Intelligence: Student / Teacher / Orchestrator

Intelligence isn't monolithic. It's layered:

### LEVEL 3: ORCHESTRATOR

- Manages multiple world branches (Fork tree)
- Preserves failed worlds (no deletion)
- Selects successful path
- NO intelligence required (deterministic)




- All proposals verified through execution
- Falsifiable and auditable

Teacher = "A modeler who can be wrong"

## Trust Boundaries

Not everything is trusted equally:



### The Firewall Principle:

LLM → Proposal → Authority → Approval → Core.apply()

NOT:

LLM → State // FORBIDDEN

LLM outputs **never** directly mutate state. Ever.

## The Runtime Flow

Here's how execution actually works:



1. PROJECTION renders Snapshot → View  
(Full state filtered to what Actor needs)
2. ACTOR (Student) observes View  
(Sees only available actions)
3. ACTOR proposes ChangeSet  
(Selects action, does NOT execute)
4. Wrap ChangeSet → Proposal  
(Attach identity for accountability)
5. AUTHORITY decides  
→ approved / rejected / changes\_requested
6. If approved, ORCHESTRATOR forwards to Core
7. CORE.executeAction()
  - Check availability (again!)
  - Resolve effect parameters
  - Execute handler → Patch[]
  - Apply patches → New Snapshot
  - Recompute affected values
8. Loop

## Failure & Fork

When things go wrong:

1. Core returns { ok: false, reason: 'UNAVAILABLE' }
2. Orchestrator captures failure
  - Preserve current Snapshot (no deletion!)
  - Generate ExplainGraph
3. Teacher (if present) analyzes
  - Input: Trace + ExplainGraph
  - Output: World Hypothesis
4. Authority decides on Fork  
→ approve / reject
5. Orchestrator creates Fork
  - Parent Runtime preserved
  - Child Runtime with hypothesis applied
6. Execution continues in Child



7. On success, Authority selects winner

Failed worlds are **preserved**, not deleted. You can always go back.

## Why This Works

### Intelligence Redistribution

BEFORE (Intelligence-Centric):

Intelligence = World Understanding  
 + Rule Enforcement  
 + Action Selection  
 + State Management  
 + Failure Diagnosis

AFTER (World-Centric):

Intelligence = Action Selection (minimal)


World Understanding → Projection  
 Rule Enforcement → Core (availability)  
 State Management → Core (Patch/Apply)  
 Failure Diagnosis → Core (Explain)

We moved 80% of the work **out of the model**.

## Model Size Independence

Intelligence-Centric:


Stability



(Bigger model = More stable)

World-Centric:

Stability



A random selector and a **smart-enough model** (e.g., GPT-5 or Claude Opus) both produce **valid** execution traces—because validity is enforced by the World (availability gates + Patch/Apply), not by the model's reasoning.

A stronger model tends to be more **efficient**, but **correctness does not depend on model capability**.

## When NOT to Use This

World-Centric isn't for everything:

Scenario	Why It Might Not Fit
Pure creative writing	No "correct" state
Open-ended chat	No action structure
Real-time streaming	Snapshot overhead
Trivial single-turn QA	Overkill

This architecture is designed for **structured, multi-step, verifiable tasks**.

## Key Takeaways

1. **Agent failures are often world modeling failures**, not intelligence failures
2. **Make the World explicit**: Immutable snapshots, declarative patches, computed availability
3. **Intelligence proposes, never executes**: The firewall principle

 4. **Stability is achievable without scaling:** A random selector can produce valid execution

 5. **Everything should explain itself:** If it can't answer "Why?", it's a black box

 6. **Preserve failures:** Fork trees enable recovery and exploration

---

## The Philosophy

"A system that cannot explain 'Why' is a dead system."

Most agent architectures are black boxes. You pump in tokens, hope for the best, and debug through prayer.

World-Centric Architecture is a **white box**. Every value has provenance. Every action has availability. Every failure has explanation.

That's not just good engineering. That's the difference between a demo and a production system.

---

## See It In Action

**Theory is nice. Working software is better.**

**TaskFlow** is a *playable proof-of-concept* built **before** the full 7-layer architecture—designed to demonstrate a simpler but critical idea:

**LLMs can be reduced to intent interpreters and semantic carriers,**

**while the system itself remains fully deterministic.**

TaskFlow shows that even *without* the complete World-Centric runtime, a deterministic machine driven by explicit state and rules can already work.

👉 Try the demo: <https://taskflow.manifesto-ai.dev>



## What's Next?



This architecture is implemented in **Manifesto**, an open-source framework. Future posts will cover:

- Deep dive into Core's 7 layers
- Building custom Actors with LLMs
- Authority patterns for different governance needs
- Fork strategies for exploration

Questions? Disagreements? Let me know in the comments.

*Thanks for reading. If you're tired of debugging agent failures at 3 AM, maybe it's time to stop blaming the model and start designing the world.*

## Top comments (0) ⚡

[Code of Conduct](#) • [Report abuse](#)



Postmark

PROMOTED

...



ActiveCampaign > **Postmark**

Your current email provider vs. Postmark

[Get started free](#)

## We know this comparison looks a bit dramatic, but...

When your password reset emails take 10 minutes to arrive and your users are refreshing their inbox like it's a broken webpage, maybe it's time for an upgrade. Postmark makes email delivery feel less like fighting ancient technology and more like, well, magic.

[Get started free](#)



Jung Sungwoo

UX/UI Software Engineer

LOCATION

Korea, Seoul

JOINED

2024년 9월 4일

+

...



## More from Jung Sungwoo

The Mind Protocol: Why Your AI Agent Needs a World Before It Can Think

#ai #architecture #typescript #agents

Buttons Are Not Functions

#ai #frontend #architecture #agents

Why Your AI Agent Says "Done" But Nothing Actually Happened

#agents #ai #automation #llm



Postmark PROMOTED

ActiveCampaign ➤  
**Postmark**



+

...

## Your current email provider vs. Postmark

[Get started free](#)

### We know this comparison looks a bit dramatic, but...

When your password reset emails take 10 minutes to arrive and your users are refreshing their inbox like it's a broken webpage, maybe it's time for an upgrade.

Postmark makes email delivery feel less like fighting ancient technology and more like, well, magic.

[Get started free](#)